

1. Model-Based Learning Agent
2. Water Jug Problem
3. Best First Search
4. Machinery (Puzzles)
5. Components of Problem
6. Resolution (Pro) and Unification
7. Inference Rule
8. Mean-End Analysis
9. Total Partial Order
10. Plan Graph
11. Rule-Based Expert System
12. Minimax Algorithm
13. Vacuum Cleaner Problem
14. Evolution of AI
15. Properties of Tark Environment
16. Hill Climbing
17. Simulated Annealing
18. Uninformed Search Techniques
19. Semantic Field
20. Rule-Based Expert System
21. Conditional Planning
22. Goal Stack Planning
23. Alpha-Beta Pruning
24. Reflex Agent

Reflex Agent in AI

Definition:

A **Reflex Agent** is an agent that chooses its actions based on the current percept, ignoring the rest of the percept history. These agents operate on a simple **condition-action rule**: if a condition is true, then an action is performed.

Reflex agents are the simplest type of agent, often suitable for environments where actions can be decided without needing to consider previous states or future consequences.

Components of a Reflex Agent

1. **Percept**: Information received from the environment at the current time.
 2. **Condition-Action Rules**: Predefined rules that map conditions (percepts) to actions.
 3. **Action**: The action chosen based on the percept and corresponding rule.
-

Types of Reflex Agents

1. **Simple Reflex Agent**:
 - Acts only on the current percept.
 - Does not maintain any internal state.
 - Useful in static and fully observable environments.
 - Example: A thermostat that turns on or off based on the temperature.
 2. **Model-Based Reflex Agent**:
 - Uses a model of the world to handle partially observable environments.
 - Keeps track of the state based on the percept history.
 - Example: A robot vacuum that avoids obstacles based on prior knowledge of the room layout.
-

Working of a Simple Reflex Agent

1. **Input**: The agent perceives its environment using sensors.
 2. **Condition Check**: The agent checks if any predefined condition matches the current percept.
 3. **Action Execution**: The agent performs the action associated with the matching condition.
-

Advantages

1. **Simplicity**: Easy to design and implement.
 2. **Fast Response**: Acts quickly since it doesn't involve complex computations or reasoning.
 3. **Low Computational Overhead**: Doesn't need to store or process percept history.
-

Disadvantages

1. **No Memory**: Cannot handle partially observable environments effectively since it doesn't store past percepts.
 2. **Limited Intelligence**: Can only handle predefined scenarios; cannot adapt or learn.
 3. **Fails in Complex Environments**: Not suitable for dynamic, unpredictable, or multi-step environments.
-

Example: Vacuum Cleaner Problem

A reflex agent for a vacuum cleaner might work as follows:

Problem Description:

The vacuum cleaner operates in an environment with two rooms: Room A and Room B. Each room can be clean or dirty.

Rules:

1. If the current room is dirty → **Clean it.**
2. If the current room is clean → **Move to the next room.**

Implementation:

- **Percepts:** [Room, Dirt status]
 - Example: [A, Dirty]
- **Actions:**
 - Clean the room.
 - Move to the next room.

Pseudocode for Simple Reflex Agent

python

Copy code

```
def reflex_vacuum_agent(percept):
    room, status = percept # Current room and dirt status

    if status == 'Dirty':
        return 'Clean'
    elif room == 'A':
        return 'Move to B'
    elif room == 'B':
        return 'Move to A'
```

Flowchart for Reflex Agent

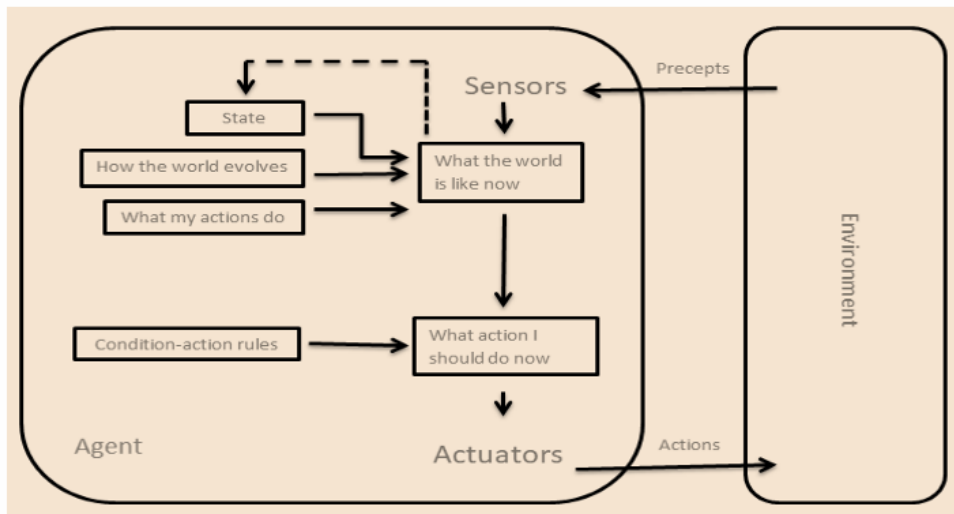
1. **Input Percept** →
2. **Match Condition** →
3. **Perform Action** →
4. **Repeat**

Applications of Reflex Agents

- **Industrial Machines:** Simple automation systems like conveyor belts.
- **Home Appliances:** Thermostats, automatic lights, and robotic vacuum cleaners.
- **Traffic Systems:** Basic traffic light controllers.

Model-Based Reflex Agents

A model-based agent can handle **partially observable environments** by the use of a model about the world.



Model-Based Learning Agent

A **Model-Based Learning Agent** is an intelligent agent that uses an internal model of the environment to reason about and predict the effects of its actions. This type of agent is more sophisticated than simple reflex agents as it can handle partially observable environments by maintaining an internal state.

Definition

- A **model-based agent** keeps track of the state of the world by maintaining a model (representation) of how the environment evolves over time and how the agent's actions affect the environment.
- The model is used to **update its belief about the state** and **select actions** that maximize its performance.

Components of a Model-Based Learning Agent

1. **Percept:**
 - The information received from the environment through sensors.
 2. **Model:**
 - An internal representation of the environment that allows the agent to infer the state of the world based on percepts and actions.
 3. **Internal State:**
 - Maintains a representation of the agent's knowledge about the world, which is updated based on the model and percepts.
 4. **Reasoning Mechanism:**
 - Uses the model and current state to predict outcomes of actions and select the optimal one.
 5. **Action:**
 - Chooses the next action based on the reasoning mechanism and the agent's goals.
-

How Model-Based Agents Work

1. **Perception:**
 - The agent observes the current state of the environment through its sensors.
 2. **Model Update:**
 - Updates its internal model of the environment based on new observations and past actions.
 3. **Reasoning:**
 - Predicts the effect of actions using the model.
 - Plans actions to move toward its goal.
 4. **Action Execution:**
 - Executes the chosen action and waits for new percepts.
-

Advantages of Model-Based Learning Agents

1. **Handles Partial Observability:**
 - Maintains an internal state to infer unobservable aspects of the environment.
 2. **Predictive Capability:**
 - Can reason about the outcomes of actions before executing them.
 3. **Flexibility:**
 - Can adapt to changes in the environment by updating its model.
 4. **Efficiency:**
 - Reduces the number of actions required to achieve a goal by planning effectively.
-

Disadvantages

1. **Complexity:**
 - Maintaining and updating the model can be computationally expensive.
 2. **Model Dependence:**
 - Performance depends on the accuracy of the model; an inaccurate model leads to suboptimal actions.
 3. **High Memory Requirement:**
 - Requires storage for maintaining the model and internal state.
-

Example: Vacuum Cleaner Problem

Problem:

A vacuum cleaner operates in an environment with two rooms: A and B. Each room can be dirty or clean, and the vacuum cleaner can sense dirt only in its current room. It needs to clean both rooms optimally.

Model-Based Learning Agent Design:

1. **Internal Model:**
 - Keeps track of which rooms are clean and which are dirty.
 2. **Percepts:**
 - [Current room, Dirt status].
 3. **Model Update:**
 - If the vacuum senses a room is dirty, mark it as dirty in the model.
 - If the vacuum cleans a room, update the model to mark the room as clean.
 4. **Action:**
 - If the current room is dirty, clean it.
 - If the current room is clean, move to the other room.
-

Algorithm for Vacuum Cleaner:

python

```
class VacuumAgent:
    def __init__(self):
        self.model = {"A": None, "B": None} # State of rooms: None means unknown

    def update_model(self, room, status):
        self.model[room] = status # Update the room status in the model

    def choose_action(self, room):
        if self.model[room] == "Dirty":
            return f"Clean {room}"
        elif room == "A":
            return "Move to B"
        elif room == "B":
            return "Move to A"

# Example Usage
agent = VacuumAgent()
agent.update_model("A", "Dirty")
print(agent.choose_action("A")) # Output: Clean A
agent.update_model("A", "Clean")
print(agent.choose_action("A")) # Output: Move to B
```

Key Applications

1. **Robotics:**
 - Robots that maintain an internal map of their surroundings.
2. **Game AI:**
 - Agents in games that reason about opponents' actions and adjust strategies dynamically.
3. **Autonomous Vehicles:**
 - Self-driving cars that model the environment, including other vehicles and pedestrians.
4. **Healthcare:**
 - Intelligent diagnostic systems that infer the progression of diseases based on symptoms and medical history.

Vacuum Cleaner Problem

The **Vacuum Cleaner Problem** is a classical problem in artificial intelligence, commonly used to demonstrate state space search, heuristic search, and agent-based problem-solving. The task is to design a vacuum cleaner agent that can clean a two-room environment (or more) and return to a "goal" state.

Problem Setup:

- **Environment:** The environment consists of a series of rooms, each of which can either be clean or dirty.
 - **Agent:** The agent (vacuum cleaner) can move between rooms, cleaning dirty rooms and returning to a designated goal state.
 - **Objective:** The agent must clean all the dirty rooms and, if needed, return to the goal state.
-

State Space Representation:

- A state can be represented as a tuple where each component represents the status of a room (clean or dirty).
 - For example, in a two-room setup:
 - State: (Clean, Dirty) means the first room is clean, and the second is dirty.
 - The agent can either clean a dirty room or move to another room.
-

Possible Actions:

1. **Move:** The vacuum cleaner can move from one room to another.
2. **Clean:** The vacuum cleaner can clean a dirty room.

These actions lead to different state transitions, depending on the environment.

Example Solution:

1. **Initial State:** (Dirty, Dirty) (Both rooms are dirty).
2. **Action 1:** Clean the first room → (Clean, Dirty).
3. **Action 2:** Move to the second room → (Clean, Dirty) (still).
4. **Action 3:** Clean the second room → (Clean, Clean).


Once all rooms are clean, the goal is achieved.

Agent's Strategy:

The vacuum cleaner agent must employ a strategy to perform actions like cleaning and moving systematically until all rooms are clean. A simple approach could be:

- If the current room is dirty, clean it.
- If the current room is clean, move to the next room.
- Repeat the process until all rooms are clean.

STATE SPACE FOR VACUUM CLEANER AGENT PROBLEM



The Water Jug Problem is a classic problem in artificial intelligence and problem-solving, typically solved using a search algorithm like Breadth-First Search (BFS) or Depth-First Search (DFS).

The Water Jug Problem Solution (4-Liter Jug and 3-Liter Jug)

In this scenario, you have two jugs:

- **Jug 1 (4-Liter Jug):** Can hold a maximum of 4 gallons.
- **Jug 2 (3-Liter Jug):** Can hold a maximum of 3 gallons.

You need to get exactly 2 gallons of water into the **4-liter jug**.

State Space Representation


Each state can be represented as a pair (x,y) , where:

- x is the amount of water in the 4-liter jug.
- y is the amount of water in the 3-liter jug.

The valid states for the problem are:

- $x \in \{0,1,2,3,4\}$ (capacity of 4-Liter jug).
- $y \in \{0,1,2,3\}$ (capacity of 3-Liter jug).

BFS for a water jug problem



Initial and Goal State

- **Initial state:** $(0,0)$ meaning both jugs are empty.
- **Goal state:** $(2,n)$ where n can be any value from 0 to 3 (you need 2 gallons in the 4-liter jug).

Possible Goal States

In addition to the main goal state of $(2,n)$, the problem also allows several possible goal states for the 3-liter jug, depending on how the operations are performed. These include:

- **(2, 0):** 2 gallons in the 4-liter jug and 0 gallons in the 3-liter jug.
- **(2, 1):** 2 gallons in the 4-liter jug and 1 gallon in the 3-liter jug.
- **(2, 2):** 2 gallons in the 4-liter jug and 2 gallons in the 3-liter jug.
- **(2, 3):** 2 gallons in the 4-liter jug and 3 gallons in the 3-liter jug.

Therefore, any state where the 4-liter jug contains 2 gallons and the 3-liter jug holds between 0 and 3 gallons is a valid goal state.

Operations/Production Rules

The following operations or production rules can be applied:

1. **Fill 4-liter jug:** $(x, y) \rightarrow (4, y)$ if $x < 4$.
2. **Fill 3-liter jug:** $(x, y) \rightarrow (x, 3)$ if $y < 3$.
3. **Empty 4-liter jug:** $(x, y) \rightarrow (0, y)$ if $x > 0$.
4. **Empty 3-liter jug:** $(x, y) \rightarrow (x, 0)$ if $y > 0$.
5. **Pour from 4-liter jug to 3-liter jug:** $(x, y) \rightarrow (x - d, y + d)$, where d is the amount of water that can be poured (up to the capacity of the 3-liter jug).
6. **Pour from 3-liter jug to 4-liter jug:** $(x, y) \rightarrow (x + d, y - d)$, where d is the amount of water that can be poured (up to the capacity of the 4-liter jug).

Solution (Step-by-Step)

Solution (Step-by-Step)

To get exactly 2 gallons in the 4-liter jug, we can follow these steps using the operations above:

1. Start at $(0, 0)$: Both jugs are empty.
2. Fill the 3-liter jug: $(0, 0) \rightarrow (0, 3)$.
3. Pour from 3-liter jug to 4-liter jug: $(0, 3) \rightarrow (3, 0)$ (Now 4-liter jug has 3 gallons).
4. Fill the 3-liter jug again: $(3, 0) \rightarrow (3, 3)$.
5. Pour from 3-liter jug to 4-liter jug: $(3, 3) \rightarrow (4, 2)$ (Now 4-liter jug has 4 gallons and 3-liter jug has 2 gallons).
6. Empty the 4-liter jug: $(4, 2) \rightarrow (0, 2)$.
7. Transfer 2 liters from 3-liter jug to 4-liter jug: $(0, 2) \rightarrow (2, 0)$.

Now, the 4-liter jug contains exactly 2 gallons, and the goal has been achieved.

Sequence of Actions:

Sequence of Actions:

- $(0, 0) \rightarrow (0, 3)$ (Fill 3-liter jug)
- $(0, 3) \rightarrow (3, 0)$ (Pour from 3-liter jug to 4-liter jug)
- $(3, 0) \rightarrow (3, 3)$ (Fill 3-liter jug again)
- $(3, 3) \rightarrow (4, 2)$ (Pour from 3-liter jug to 4-liter jug)
- $(4, 2) \rightarrow (0, 2)$ (Empty 4-liter jug)
- $(0, 2) \rightarrow (2, 0)$ (Transfer 2 liters from 3-liter jug to 4-liter jug)


State Space Search

The state space search involves generating all possible states by applying the production rules and checking if the goal state is reached. The search involves exploring the different combinations of water in the two jugs until we find a state where the 4-liter jug contains exactly 2 gallons.

Alternative Solutions

There may be alternative solutions or different sequences of actions to reach the goal. For example, in a different sequence, you may perform other valid moves, and you may reach the goal by different transitions.

DFS for a water jug problem



Production Systems

A **production system** is a model used in artificial intelligence (AI) to **describe a system that generates or produces rules (states)** to reach a solution. It is a framework for solving problems by applying rules to states in a defined order until a goal is reached.

A typical production system consists of four basic components:

1. Set of Rules (Ci, Ai)

- **Ci (Condition):** Describes the current state or condition.
- **Ai (Action):** Defines the action to be taken if the condition is met.
- Rules transform one state to another.

2. Knowledge Databases

- Stores relevant facts and information about the problem.
- Includes data like problem parameters, initial states, and possible actions.

3. Control Strategy

- Determines the **order** in which rules are applied.
- Uses strategies like breadth-first or depth-first search to find the goal.

4. Rule Applier

- Applies rules to the current state based on the control strategy.
- Transitions the system from one state to another until the goal is reached.

Components of a Problem in AI

When designing or solving a problem in Artificial Intelligence, it is essential to represent the problem clearly and break it into specific components. The **components of a problem** define how an AI agent perceives and acts to achieve a solution.

Components of a Problem

1. Initial State

- The starting point or configuration of the system when the problem-solving process begins.
- Example: For the **8-puzzle problem**, the initial state is the starting arrangement of tiles.

2. Goal State

- The desired state that represents the solution to the problem.
- Example: For the **8-puzzle problem**, the goal state is when all tiles are in numerical order.

3. State Space

- The set of all possible states the system can be in, including the initial state, intermediate states, and goal state.
- Example: In the **8-puzzle problem**, the state space consists of all possible tile arrangements.

4. Actions (Operators)

- The set of legal actions available to the agent that can transition the system from one state to another.
- Example: For the **8-puzzle problem**, actions include moving tiles left, right, up, or down.

5. Transition Model

- Defines how the actions change the current state to a new state. It's a function describing the result of applying an action to a state.
- Example: If you move the blank tile up in the **8-puzzle**, the transition model defines the new arrangement of tiles.

6. Path Cost

- The cost associated with a sequence of actions taken to reach the goal state. This helps in finding the optimal solution.
- Example: The path cost may represent the number of moves in the **8-puzzle** or the distance traveled in a navigation problem.

1. Introduction

A **Rule-Based Expert System** (RBES) is an AI system that uses a knowledge base of "if-then" rules to solve complex problems within a specific domain. These systems are designed to mimic the decision-making abilities of human experts by applying logical rules to the given data.

2. Overview of Rule-Based Expert System

Rule-based expert systems use knowledge in the form of rules to reason through problems and make decisions. The system's core components include:

- **Knowledge Base:** Contains the set of rules and facts.
- **Inference Engine:** The mechanism that applies the rules to the known facts to infer conclusions.
- **User Interface:** Allows interaction between the user and the system to provide input and receive conclusions.

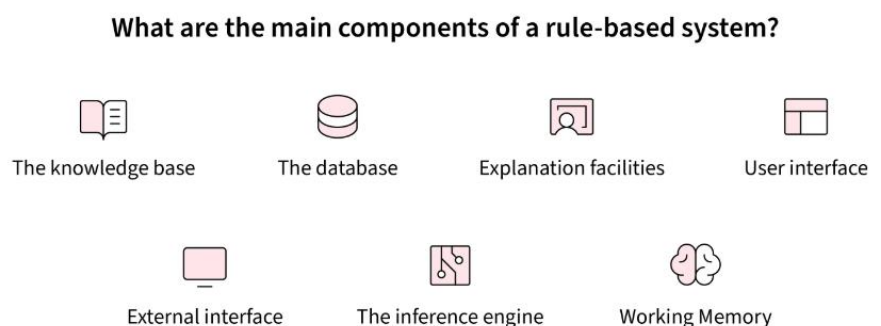

3. Evolution of Rule-Based Expert System

The concept of rule-based systems began in the 1970s when researchers tried to simulate expert human reasoning through symbolic logic and programming. Early rule-based systems, like **MYCIN**, were designed for medical diagnosis. Over time, advancements in AI and computing led to more sophisticated rule-based systems, incorporating advanced reasoning mechanisms and more complex rules.

4. Components of Rule-Based Expert System

The main components of a Rule-Based Expert System are:

- **Knowledge Base:** Contains domain-specific knowledge, often encoded as rules in the form of **if-then** statements.
- **Inference Engine:** The reasoning mechanism that interprets the rules and applies them to known facts to derive conclusions.
- **Working Memory:** A dynamic part of the system where the current facts (known data) are stored.
- **User Interface:** Allows users to interact with the system and enter queries or input data.



5. Construction of a Rule-Based Expert System

1. Create the Knowledge Base:

- The first step is to create the Knowledge Base, where rules (if-then statements) and facts are stored. These rules define how the system should make decisions based on the facts it receives.

2. Develop the Inference Engine:

- Next, design the Inference Engine, which is responsible for applying the rules to the facts and using logical reasoning to infer conclusions. The engine can use forward or backward chaining to derive solutions.

3. Integrate External Databases:

- After the basic components are set, integrate External Databases to provide additional facts or information not included in the internal knowledge base. This enhances the system's ability to consider more variables.

4. Incorporate External Programs:

- The system may also include External Programs that add supplementary functionalities, enabling it to perform tasks or solve complex problems that the core RBES cannot handle on its own.

5. Enable Explanation Facilities:

- Explanation Facilities are incorporated to allow the system to explain its reasoning process. This helps users understand how conclusions are reached, improving transparency and trust in the system.

6. Develop the User Interface:

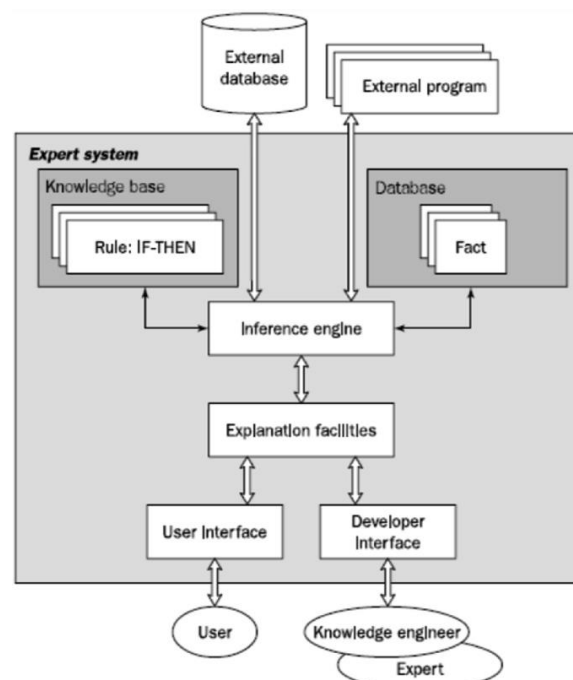
- A User Interface is created, allowing users to input their data or queries and receive responses from the system. The interface must be intuitive to ensure ease of use.

7. Create the Developer Interface:

- Alongside the user interface, a Developer Interface is built for knowledge engineers to maintain and update the system. They can add or modify the knowledge base and rules to keep the system relevant.

8. Define Roles of User, Knowledge Engineer, and Expert:

- **User:** Interacts with the system to obtain solutions or recommendations.
- **Knowledge Engineer:** Updates and maintains the knowledge base, ensuring its accuracy and relevance.
- **Expert:** Provides domain-specific knowledge to populate the knowledge base and ensures it is accurate.



6. Pros and Cons of Rule-Based Expert Systems

Pros:

- **Simplicity:** Easy to understand and implement.
- **Transparency:** Clear reasoning process through traceable rules.
- **Efficiency:** Works well in defined, structured domains.

Cons:

- **Scalability:** Managing and maintaining a large set of rules can be difficult.
 - **Lack of Flexibility:** Struggles in dynamic or unstructured environments.
 - **Dependence on Rule Quality:** Inaccurate or incomplete rules lead to unreliable outputs.
-

7. Applications of Rule-Based Expert Systems

RBES is used in various fields, such as:

- **Medical Diagnosis:** MYCIN diagnoses diseases based on symptoms.
- **Customer Support:** Automates troubleshooting based on predefined solutions.
- **Financial Advisory:** Assesses loan eligibility and detects fraud in banking.

Legal Systems: Assists in analyzing case laws and generating legal advice.

Properties of Task Environment in AI

The **task environment** is the "world" in which an agent operates and makes decisions to achieve its goals. Understanding the properties of a task environment is critical for designing an intelligent agent. These properties define the complexity of the environment and determine the type of AI techniques that should be used for the agent.

Key Properties of a Task Environment

1. Fully Observable vs. Partially Observable

- **Fully Observable:** The agent has access to the complete state of the environment at any given time.
 - Example: Chess (the entire board is visible).
 - **Partially Observable:** The agent has incomplete or noisy information about the environment.
 - Example: Poker (you cannot see other players' cards).
-

2. Deterministic vs. Stochastic

- **Deterministic:** The next state of the environment is completely determined by the current state and the agent's action.
 - Example: Tic-Tac-Toe.
 - **Stochastic:** The next state is uncertain and involves randomness.
 - Example: Rolling a die in a board game.
-

3. Episodic vs. Sequential

- **Episodic:** Each decision (episode) is independent of the previous one. The agent does not need to remember past actions.
 - Example: Image classification (each image is processed independently).
 - **Sequential:** Current actions affect future actions and outcomes. The agent must consider the history of its actions.
 - Example: Self-driving cars.
-

4. Static vs. Dynamic

- **Static:** The environment does not change while the agent is deliberating.
 - Example: Crossword puzzles.
 - **Dynamic:** The environment changes over time, even as the agent is deciding.
 - Example: Traffic while driving.
-

5. Discrete vs. Continuous

- **Discrete:** The environment has a finite number of states, actions, and outcomes.
 - Example: Chess (finite board configurations).
 - **Continuous:** The environment has an infinite number of possible states and actions.
 - Example: Robot navigation in a real-world space.
-

6. Single-Agent vs. Multi-Agent

- **Single-Agent:** The agent interacts with the environment alone.
 - Example: Solving a maze.

- **Multi-Agent:** The agent interacts with other agents, which may be cooperative, competitive, or neutral.
 - Example: Soccer game simulation (multiple players).
-

7. Known vs. Unknown

- **Known:** The agent has prior knowledge about the environment, including rules and dynamics.
 - Example: Solving a Sudoku puzzle.
 - **Unknown:** The agent has to learn about the environment by interacting with it.
 - Example: Exploring an unknown maze.
-

Example: Properties of the Vacuum Cleaner Problem

Let's analyze the **vacuum cleaner problem** using these properties:

1. **Observability:** Partially Observable (the vacuum can sense dirt only in the current location).
 2. **Deterministic/Stochastic:** Deterministic (cleaning a location always removes dirt).
 3. **Episodic/Sequential:** Sequential (the vacuum must consider previous actions to clean efficiently).
 4. **Static/Dynamic:** Static (the environment does not change unless the vacuum acts).
 5. **Discrete/Continuous:** Discrete (finite locations: A and B).
 6. **Single-Agent/Multi-Agent:** Single-Agent (only the vacuum interacts with the environment).
 7. **Known/Unknown:** Known (the vacuum knows the rules of the environment).
-

Importance of Understanding Task Environment

- **Designing the Agent:** Helps in selecting the most suitable algorithm or learning method for the agent.
- **Efficiency:** Knowing properties like determinism and observability aids in optimizing performance.
- **Scalability:** Helps in preparing agents for more complex environments.

Simulated Annealing: An AI Optimization Technique

Simulated Annealing (SA) is a probabilistic algorithm inspired by the physical process of annealing in metallurgy, where metals are heated and then gradually cooled to reach a stable, low-energy state. In AI, SA is used for optimization problems, offering a balance between completeness and efficiency by avoiding local optima.

Key Concepts:

1. **Hill-Climbing Limitation:** Traditional hill-climbing algorithms get stuck in local maxima, which is not ideal for global optimization.
2. **Simulated Annealing Advantage:** Unlike hill-climbing, SA accepts worse solutions with a certain probability, allowing it to escape local maxima and potentially find a global optimum.

Process of Simulated Annealing:

1. **Initialization:** Start with a random solution and define a temperature schedule.
2. **Iteration:**

- Randomly select a neighboring solution.
- If the solution improves, accept it. If not, accept it with a probability $e^{\frac{-\Delta E}{T}}$ (where T is the temperature).

3. **Cooling:** Gradually decrease the temperature to focus more on refining the solution.
4. **Termination:** The algorithm stops when the temperature reaches zero or another stopping criterion is met.

Pseudocode:

python

```
def simulated_annealing(problem, schedule):
    current = make_node(initial_state(problem))
    best = current

    for t in range(1, infinity):
        T = schedule[t]
        if T == 0:
            return current

        next = randomly_select_successor(current)
        deltaE = value(next) - value(current)

        if deltaE > 0 or random.uniform(0, 1) < math.exp(-deltaE / T):
            current = next

        if value(current) < value(best):
            best = current

        T *= cooling_rate

    return best
```

Applications:

- **Optimization Problems:** Used in problems like the **Traveling Salesman Problem (TSP)**, **neural network training**, and **hyperparameter tuning**.
 - **Resource Scheduling:** Can be applied to job scheduling and circuit design.
-

Advantages:

- **Global Optimization:** Avoids local optima and can explore a broader solution space.
 - **Flexibility:** Works for various types of optimization problems (both continuous and discrete).
 - **Simplicity:** Easy to implement and doesn't require complex mathematical tools.
-

Disadvantages:

- **Slow Convergence:** Can be computationally expensive for large problems.
- **Sensitive to Parameters:** Performance depends heavily on temperature schedules and cooling rates.

Conditional Planning in AI

Conditional Planning is a planning technique used when the agent operates in environments where uncertainty exists, making it difficult to predict the outcome of actions or even the initial state of the environment. In such environments, plans need to be adaptable, taking into account different possible outcomes and contingencies. This is in contrast to traditional planning, where the agent executes a predefined set of actions.

Key Concepts of Conditional Planning:

1. Uncertainty in the Environment:

- The agent may not always know its current state or the results of its actions. For example, in a partially observable environment, an agent may not be aware of all aspects of its surroundings.

2. Branches in Plans:

- In conditional planning, instead of a straight-line sequence of actions, plans include branches that depend on conditions being true or false.
- A typical structure is:
 - **if <condition> then <plan A> else <plan B>**

3. Full Observability:

- **Full observability** means that the agent knows exactly what the current state is without needing to perform an observation action. This ensures the agent can plan for the most accurate possible situation.
- The agent can plan for all possible contingencies, considering multiple scenarios that might unfold.

4. Plans for Contingencies:

- The agent needs to prepare plans that account for various possible contingencies. This might involve having alternative plans if the primary plan fails due to unexpected results from actions.

Example: The Vacuum World Problem (Double Murphy World)

In the **vacuum world**, the agent is tasked with cleaning a 2x2 grid. The agent can move left or right and can apply a "suck" action to clean dirt. However, the environment is uncertain because sometimes:

- The vacuum cleaner might deposit dirt when it moves to a clean square.
- The vacuum cleaner might also deposit dirt if the "suck" action is applied to a clean square.

This creates an environment where the agent needs to plan for these uncertain outcomes, considering multiple possibilities.

Plan with Conditional Branches in the Double-Murphy Vacuum World:

- The agent's plan involves **conditional steps**, depending on whether certain conditions hold true (like whether the square is clean or dirty).

Here is an example of a **conditional plan** in the Double Murphy vacuum world:

Plan:

- **Move Left.**
- **If** the agent is at the left square, and both squares are clean ($AtL \wedge CleanL \wedge CleanR$), **then do nothing.**
- **Else** apply the **Suck** action to clean the left square.

This plan uses the conditional approach where the agent reacts to the environment's state by choosing different actions based on whether the squares are clean or dirty.

Performing and-Or Search:

And-Or Search is often used in conditional planning to handle the search for solutions that involve multiple branches (or choices) at different points. In an **And-Or Search**:

- The "**And**" represents the requirement to achieve multiple subgoals simultaneously.
- The "**Or**" represents a choice between multiple alternative actions, depending on the conditions at the time.

In the context of the vacuum world, the agent must decide between various alternatives (such as whether to suck or not, based on the cleanliness of the squares).

Example Search:

- **Goal:** Clean the vacuum world.
- The agent performs a sequence of actions:
 - **Left, Suck**
 - **Right, Suck**
 - The plan depends on the conditions met after each action.

Summary of Key Points:

- **Conditional Planning** is useful in uncertain environments where the outcomes of actions are not predictable.
- Plans contain **branches** that adapt to various conditions.
- In the **vacuum world**, the agent may face uncertainty about the cleanliness of squares or the side effects of its actions.
- **Full observability** ensures that the agent has enough information to act intelligently.
- **And-Or search** helps the agent consider multiple possibilities and plan for various contingencies.

Conditional planning enables intelligent decision-making in environments where outcomes are uncertain, helping the agent to adapt and respond to changing conditions.

40 mini

Means-Ends Analysis (MEA) Overview

Means-Ends Analysis (MEA) is a **problem-solving strategy** in Artificial Intelligence (AI) that combines forward and backward reasoning to solve complex problems. It breaks down a problem into smaller subproblems (subgoals) by identifying differences between the current state and the goal state. MEA then applies operators to reduce those differences until the goal is achieved.

Key Features of Means-Ends Analysis

1. **Mixed Search Strategy:** MEA combines forward and backward search techniques to solve problems efficiently.
 2. **Problem Decomposition:** It divides a large problem into smaller subproblems, solving them step by step.
 3. **Recursive Process:** MEA applies itself recursively, working on differences between states until the goal is reached.
 4. **Goal-Driven:** MEA focuses on achieving the goal state by iteratively reducing the gap from the current state.
-

How Means-Ends Analysis Works

MEA involves a series of steps, which are executed recursively:

1. **Evaluate the Difference:** Compare the **current state** with the **goal state** and identify differences.
 2. **Select an Operator:** Choose an operator (action) that can reduce the difference. If no operator exists, signal failure.
 3. **Apply the Operator:**
 - Define two intermediate states:
 - **O-Start:** A state where the operator's preconditions are satisfied.
 - **O-Result:** The state after applying the operator.
 4. **Recursive MEA Calls:**
 - Solve the subproblem of reaching **O-Start** from the current state.
 - Solve the subproblem of reaching the goal state from **O-Result**.
 5. **Combine Results:** If both recursive calls are successful:
 - Combine the first part (reaching O-Start), the operator, and the last part (reaching the goal).
-

Operator Subgoal

Sometimes, the operator cannot be directly applied to the current state. In such cases, a subgoal is created to satisfy the operator's preconditions. This process, called **Operator Subgoal**, involves using backward chaining to achieve the prerequisites for the operator to function.

Algorithm for Means-Ends Analysis

plaintext

Copy code

Step 1: Compare CURRENT and GOAL.

- If there are no discrepancies, return Success and exit.

Step 2: Identify the most significant difference.

Step 3: Select an operator O that can reduce this difference.

- If no operator exists, signal failure and exit.

Step 4: Apply the operator O:

- Define O-Start (preconditions satisfied) and O-Result (state after applying O).

Step 5: Recursively apply MEA:

- Solve MEA(CURRENT, O-START).
- Solve MEA(O-RESULT, GOAL).

Step 6: If both recursive calls succeed:

- Combine results: FIRST-PART, O, LAST-PART.
- Return Success and the combined result.

Example 1: Getting a Glass of Water

Let's solve the problem using MEA.

1. **Current State:** You are in the living room without water.
2. **Goal State:** You have a glass of water.

Steps:

1. **Compare CURRENT and GOAL:** There is a discrepancy (you don't have water).
 2. **Identify the Difference:** You need water, so choose the operator "**Go to Kitchen**".
 - **O-Start:** You are in the living room (precondition for moving to the kitchen).
 - **O-Result:** You are in the kitchen (where water is available).
 3. **Recursive Calls:**
 - First Part: MEA(CURRENT, O-START) – Move from living room to kitchen.
 - Last Part: MEA(O-RESULT, GOAL) – Fill a glass with water in the kitchen.
 4. **Combine Results:** If successful:
 - Action Sequence: **Move to Kitchen** → **Fill Glass with Water**.
-

Example 2: Solving a Puzzle

Consider a problem where the initial state has a square outside a circle with a dot inside the circle, and the goal state has the square inside the circle without a dot.

Initial State:

- Square outside circle.
- Dot inside circle.

Goal State:

- Square inside circle.
- No dot.

Operators:

1. **Delete Operator:** Removes the dot.
2. **Move Operator:** Moves the square inside the circle.
3. **Expand Operator:** Adjusts the size of the square.

Steps:

1. **Step 1: Evaluate Initial State:** Compare the initial state to the goal state.
 - Difference: Dot exists in the initial state but not in the goal state.
2. **Apply Delete Operator:** Remove the dot.
 - Resulting State: Square outside the circle, no dot.
3. **Apply Move Operator:** Move the square inside the circle.
 - Resulting State: Square inside circle, no dot.
4. **Apply Expand Operator:** Adjust the size of the square.
 - Final State: Matches the goal state.

Action Sequence: Delete Dot → Move Square → Expand Square.

Advantages of Means-Ends Analysis

1. Efficient problem-solving by narrowing down the search space.
2. Handles complex problems by breaking them into smaller, manageable parts.
3. Reduces redundancy by focusing only on significant differences.
4. Combines the strengths of both forward and backward search techniques.

Disadvantages of Means-Ends Analysis

1. It may fail if there are no operators to address a difference.
 2. Recursive calls can lead to high computational costs for large problems.
 3. Relies heavily on the availability of appropriate operators.
-

Conclusion

Means-Ends Analysis is a powerful and systematic problem-solving technique in AI that bridges the gap between the current state and the goal state. By evaluating differences and recursively solving subproblems, MEA provides an efficient pathway to achieve complex goals.

Partial Order Planning vs Total Order Planning

Partial Order Planning (POP) and **Total Order Planning (TOP)** are two strategies used in Artificial Intelligence (AI) for automated planning. They differ in how they organize and order actions to achieve a goal. Let's analyze the problem of putting on a pair of shoes using these two approaches.

Problem Description

Goal:

Achieve the state where both shoes are on:

$\text{RightShoeOn} \wedge \text{LeftShoeOn}$

Operators:

1. **Op1:**
 - **Action:** RightShoe
 - **Precondition:** RightSockOn
 - **Effect:** RightShoeOn
2. **Op2:**
 - **Action:** RightSock
 - **Precondition:** None
 - **Effect:** RightSockOn
3. **Op3:**
 - **Action:** LeftShoe
 - **Precondition:** LeftSockOn
 - **Effect:** LeftShoeOn
4. **Op4:**
 - **Action:** LeftSock
 - **Precondition:** None
 - **Effect:** LeftSockOn

What is Partial Order Planning?

Partial Order Planning (POP) is a planning approach where actions are not necessarily ordered in a strict sequence. Instead, actions are partially ordered, meaning only the necessary actions are ordered to satisfy preconditions or avoid conflicts. Actions that are independent of each other can be executed in parallel.

Key Features of POP:

- Flexible
- Parallel Execution
- Less Constrained

Steps:

1. **Goal Decomposition:**

The goal is $\text{RightShoeOn} \wedge \text{LeftShoeOn}$.
This leads to two subgoals:

 - RightShoeOn
 - LeftShoeOn.
2. **Subgoal 1: Achieve RightShoeOn:**
 - Select Op1: RightShoe.
 - Precondition for Op1: RightSockOn.
 - Add Op2: RightSock to achieve RightSockOn.

3. **Subgoal 2: Achieve LeftShoeOn:**

- Select Op3: LeftShoe.
- Precondition for Op3: LeftSockOn.
- Add Op4: LeftSock to achieve LeftSockOn.

4. **Plan:**

Since the subgoals are independent, the actions can be partially ordered:

- Op2: RightSock → Op1: RightShoe
- Op4: LeftSock → Op3: LeftShoe

POP Output Plan:

1. Put on RightSock.
2. Put on LeftSock.
3. Put on RightShoe.
4. Put on LeftShoe.

Advantages of POP:

- More flexible, allowing parallel execution of independent actions (RightSock and LeftSock can be done simultaneously).
 - Reduces unnecessary constraints on action order.
-

What is Total Order Planning?

Total Order Planning (TOP) is a planning approach in which all actions are ordered in a strict sequence. Each action is dependent on the previous one, and there is no parallel execution allowed. In TOP, every step must be explicitly ordered, even if some actions could be independent.

Key Features of TOP:

- Strict Sequential Ordering
- No Parallel Execution
- Simpler

Steps:

1. **Goal Decomposition:**

Same as in POP, decompose the goal into subgoals:

- RightShoeOn
- LeftShoeOn.

2. **Subgoal Execution:**

All actions must be fully ordered. A possible sequence is:

- Op2: RightSock
- Op1: RightShoe
- Op4: LeftSock
- Op3: LeftShoe

TOP Output Plan:

1. Put on RightSock.
2. Put on RightShoe.
3. Put on LeftSock.
4. Put on LeftShoe.

Advantages of TOP:

- Simpler to implement.
- Ensures actions are executed in a fixed order.

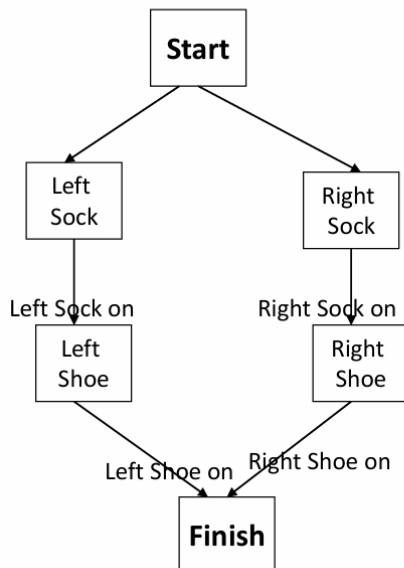
Disadvantages of TOP:

- Less flexible; cannot exploit parallelism.
- May result in unnecessary constraints and inefficiencies.

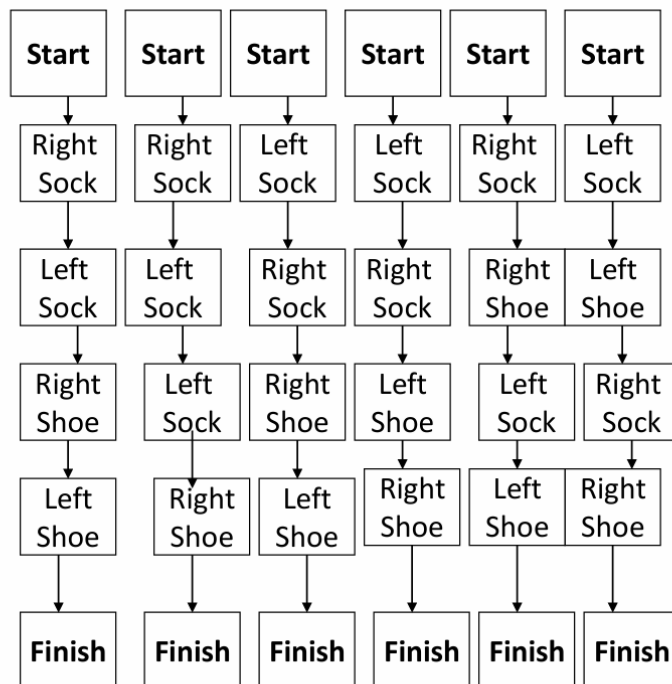
Comparison

Aspect	Partial Order Planning	Total Order Planning
Action Order	Actions are ordered only when necessary.	Actions are fully ordered in a linear sequence.
Flexibility	High; allows parallel execution.	Low; all actions are strictly ordered.
Efficiency	More efficient in problems with independent subgoals.	Less efficient due to sequential execution.
Example Output Plan	RightSock, LeftSock, RightShoe, LeftShoe (parallel where possible).	RightSock, RightShoe, LeftSock, LeftShoe

Partial Order Plans:



Total Order Plans:



Search Techniques in Artificial Intelligence

Search techniques in AI refer to strategies or algorithms used to find solutions to problems. These problems are usually defined by a set of states and actions, and the objective is to search through this space of possible states to find a solution or goal.

Search techniques are classified based on how they explore the state space, whether they use prior knowledge (heuristics), and how they handle the search process. They are commonly categorized into two main types: **uninformed (blind) search** and **informed (heuristic) search**.

Types of Search Techniques

1. Uninformed (Blind) Search:

- **Definition:** These algorithms do not have any information about the goal other than the problem definition itself. They explore the search space without any additional guidance.
- **Other names:** Blind search, Uninformed search.
- **Examples:** BFS (Breadth-First Search), DFS (Depth-First Search), Uniform Cost Search.

2. Informed (Heuristic) Search:

- **Definition:** These algorithms use domain-specific information (heuristics) to guide the search process, helping to find solutions more efficiently by evaluating the cost or distance to the goal.
- **Other names:** Heuristic search, Intelligent search, Best-First search.
- **Examples:** A* Search, Greedy Best-First Search.

1. General Search Algorithms (Uninformed Search Methods)

- **Breadth-First Search (BFS):**
 - **Description:** Explores all nodes at the present depth level before moving on to nodes at the next depth level. It guarantees the shortest path in an unweighted graph.
 - **Category:** Uninformed search.
- **Uniform Cost Search:**
 - **Description:** Expands the node with the least path cost, ensuring that the least costly solution is found. It works like BFS, but it accounts for varying edge costs.
 - **Category:** Uninformed search.
- **Depth-First Search (DFS):**
 - **Description:** Explores as far down a branch as possible before backtracking. It is memory efficient but can get stuck in deep, infinite paths.
 - **Category:** Uninformed search.
- **Depth-Limited Search (DLS):**
 - **Description:** A variation of DFS that limits the depth of exploration to avoid infinite paths, making it useful for large search spaces.
 - **Category:** Uninformed search.
- **Iterative Deepening Search (IDS):**
 - **Description:** Combines the space efficiency of DFS and the optimality of BFS. It performs a series of depth-limited searches, increasing the depth limit at each iteration.
 - **Category:** Uninformed search.

2. Informed Search Algorithms

- **Generate and Test:**

- **Description:** A simple method of search where solutions are generated and then tested to see if they satisfy the goal. It can be inefficient if the search space is large.
- **Category:** Informed search (though it can be considered a simple heuristic search).
- **Best-First Search (BFS):**
 - **Description:** Expands the most promising node based on a heuristic evaluation function. It selects nodes based on an estimate of the cost from the current node to the goal.
 - **Category:** Informed search.
- **A Search*:**
 - **Description:** An optimal and complete search algorithm that uses both the path cost and a heuristic to determine the most promising path. It is widely used in AI for pathfinding and graph traversal.
 - **Category:** Informed search.
- **Memory-Bounded Heuristic Search:**
 - **Description:** This class of search algorithms uses memory efficiently to allow for the exploration of large spaces. It limits the amount of memory used during the search process.
 - **Category:** Informed search.

3. Local Search Algorithms and Optimization Problems

- **Hill Climbing:**
 - **Description:** A local search algorithm that moves towards the direction of increasing value (i.e., it keeps moving to the neighbor with the highest value) to find the peak or goal state.
 - **Category:** Local search / Optimization search.
- **Simulated Annealing:**
 - **Description:** An optimization algorithm that mimics the cooling process of metal. It explores the search space by allowing occasional moves to worse states to avoid local optima, and the likelihood of making such moves decreases over time.
 - **Category:** Local search / Optimization search.

Summary of Techniques Categorized

Category	Techniques
Uninformed Search	BFS, Uniform Cost Search, DFS, DLS, IDS
Informed Search	Generate and Test, Best-First Search, A* Search, Memory-Bounded Heuristic Search
Local Search / Optimization	Hill Climbing, Simulated Annealing

Heuristic Functions in AI

A **heuristic function** is a key concept in AI search algorithms that provides a way to guide the search process toward the goal. It is used to evaluate the "goodness" or potential of a particular state or node in a search space.

Definition:

A **heuristic function** is a function $h(n)$ that estimates the cost or distance from a given state or node n to the goal state. The heuristic helps the algorithm decide which paths are more promising to explore next, improving search efficiency.

Purpose:

- **Guidance:** Heuristic functions help search algorithms determine the direction in which to search next by evaluating how close a node is to the goal.
- **Efficiency:** They reduce the amount of exploration needed, helping to focus on the most promising paths or states.

Key Features of Heuristic Functions:

1. **Estimation:** Heuristics estimate the cost of reaching the goal from a given node. They are not guaranteed to be exact, but they aim to give a good approximation.
2. **Domain-Specific:** Heuristics are often problem-specific. For example, in a pathfinding problem, the heuristic might be the straight-line distance (Euclidean distance) from the current node to the goal.
3. **Used in Informed Search:** Heuristic functions are used in informed search strategies like **A* search**, **Greedy Search**, and **Best-First Search** to guide the search process towards the most promising nodes.

Local Search Algorithms - Generate and Test, Heuristic Search

Generate-and-Test Algorithm

The **Generate-and-Test** algorithm is a basic approach where:

1. A potential solution is generated.
 2. The solution is tested against a goal to check if it's valid.
 3. If the solution is correct, the process stops; otherwise, a new solution is generated and tested again.
- **Features:** Simple but inefficient for large problem spaces.

Local Search Algorithms

- **Characteristics:** These algorithms operate on a single current state rather than multiple paths and move to neighboring states, not retaining the paths. They use little memory and can find solutions in large state spaces where other methods are inefficient.
 - **State Space Landscape:** This is defined by location (states) and elevation (objective function values, such as cost or profit).
-

Hill Climbing Algorithm

Hill Climbing is an iterative algorithm where the search starts at a random state, then moves to a neighboring state with the highest objective function value, repeating until no further improvements can be made.

- **Types:**
 1. **Simple Hill Climbing:** Only one neighboring node is considered at a time. It's fast but can give sub-optimal solutions.
 2. **Steepest-Ascent Hill Climbing:** Compares all neighbors and selects the best one, requiring more computation.
 3. **Stochastic Hill Climbing:** Randomly selects a neighbor and decides whether to move to it based on improvement.
 - **Problems:**
 - **Local Maximum:** A peak that's higher than its neighbors but has other higher peaks. **Solution:** Backtracking to explore different paths.
 - **Plateau:** A flat region with no discernible direction for movement. **Solution:** Taking larger steps or randomly selecting distant states.
 - **Ridges:** Sloped areas that are hard to reach in a single move. **Solution:** Using bidirectional search or moving in various directions.
-

Simulated Annealing

Simulated Annealing is an optimization algorithm inspired by the physical process of annealing in metals, where it gradually reduces temperature to reach an optimal solution.

- **Process:** It evaluates a current state, and if a neighbor provides a better value, it moves to that state. If not, it moves based on a probability function related to temperature, allowing occasional acceptance of worse solutions to escape local minima.
 - **Advantages:** Useful for finding global optima in complex problem spaces.
-

Problems and Solutions in Hill Climbing:

- **Local Maximum:** A peak that is better than neighbors but not the highest overall. **Solution:** Backtrack and explore other paths.
 - **Plateau:** Flat areas where no improvement direction is obvious. **Solution:** Take large steps or choose random distant states.
 - **Ridges:** Areas that are sloped and not reachable in a single move. **Solution:** Use bidirectional search or explore multiple directions.
-

Advantages of Hill Climbing:

- **Application:** Suitable for optimization problems like job scheduling, traveling salesperson, chip designing, and portfolio management.
- **Efficiency:** It requires minimal computation power compared to other exhaustive search algorithms.

Reasoning and Its Types

Reasoning is the process of drawing conclusions or making inferences based on available knowledge, facts, or premises. It is a fundamental aspect of artificial intelligence (AI) and logical problem-solving, where conclusions are derived systematically.

There are two main types of reasoning:

1. Forward Chaining (Data-Driven Reasoning)
2. Backward Chaining (Goal-Driven Reasoning)

1. Forward Reasoning

Forward chaining is a data-driven reasoning method that starts with known facts and applies inference rules to derive new facts until a goal or conclusion is reached.

- It is often used in **expert systems** and **real-time decision-making systems**, where all the facts need to be processed to derive new insights.
- Forward chaining is suitable when **all data is available upfront**, and the aim is to explore or discover all possible outcomes systematically.

Steps in Forward Reasoning:

1. **Initial constraints:** The system is given one or more constraints (facts or data).
2. **Rule search:** The system searches for rules in the knowledge base that match these constraints (i.e., the IF part of the rule).
3. **Apply rules:** The selected rules are applied, producing new conditions or facts from the THEN part of the rule.
4. **Repeat:** These new conditions are processed again, and the steps repeat until no new facts are generated.

Example:

- **Problem:** Prove that Colonel West is a criminal, given the following knowledge:
 - It's a crime for an American to sell weapons to hostile nations.
 - The country Nono, an enemy of America, has missiles, which were sold to them by Colonel West, an American.
- **Knowledge base:**

Example knowledge base contd.

... it is a crime for an American to sell weapons to hostile nations:

$American(x) \wedge Weapon(y) \wedge Sells(x,y,z) \wedge Hostile(z) \Rightarrow Criminal(x)$

Nono ... has some missiles, i.e., $\exists x Owns(Nono,x) \wedge Missile(x)$:

$Owns(Nono,M_1) \wedge Missile(M_1)$

... all of its missiles were sold to it by Colonel West

$Missile(x) \wedge Owns(Nono,x) \Rightarrow Sells(West,x,Nono)$

Missiles are weapons:

$Missile(x) \Rightarrow Weapon(x)$

An enemy of America counts as "hostile":

$Enemy(x,America) \Rightarrow Hostile(x)$

West, who is American ...

$American(West)$

The country Nono, an enemy of America ...

$Enemy(Nono,America)$

Simplified Forward Chaining Algorithm:

function ForwardChaining(KB, known_facts, goal):


```

while True:
    new_facts = []
    for rule in KB:
        if rule's conditions are in known_facts:
            new_fact = apply rule's conclusion
            if new_fact is not in known_facts:
                new_facts.append(new_fact)
    if new_facts is empty:
        return "Goal not reached"
    known_facts += new_facts
    if goal is in known_facts:
        return "Goal reached"

```

Example of Forward Chaining:

Scenario: Prove that *Colonel West is a criminal*.

1. **Facts:**

- Colonel West is American.
- Nono is a hostile nation.
- Missiles are weapons.

2. **Rules:**


- If an American sells weapons to a hostile nation, then they are a criminal.
- If someone is American and has sold missiles to Nono, they have sold weapons.

3. **Goal:** Determine if Colonel West is a criminal.

Process:

- Start with known facts (e.g., "Colonel West is American").
- Apply rules to derive "Colonel West sold missiles to Nono" and then "Colonel West sold weapons to a hostile nation".
- Infer the conclusion: "Colonel West is a criminal".

Forward chaining example



2. Backward Chaining

Definition:

Backward chaining is a **goal-driven reasoning** method that starts with a specific goal or hypothesis and works backward to **determine if there is enough evidence** to support the goal.

Key Characteristics:

- Begins with a **goal** or conclusion.
- Traces backward through **rules** to **determine which facts need to be true** to satisfy the goal.
- Stops when the goal is proven or when it is determined that it cannot be proven.

Steps in Backward Chaining:

1. Identify the goal or hypothesis.
2. Check if the goal can be directly proven by the known facts.
3. If not, find rules that lead to the goal and identify their premises.
4. Recursively try to prove these premises as sub-goals.
5. Continue until:
 - The goal is proven.
 - It is determined that the goal cannot be achieved.

Backward Chaining Algorithm:

plaintext

Copy code

```
function BackwardChaining(KB, goal):
    if goal is in known_facts:
        return True
    for rule in KB:
        if rule's conclusion matches the goal:
            all_conditions = rule's conditions
            if all conditions can be proven using BackwardChaining:
                return True
    return False
```

Example of Backward Chaining:

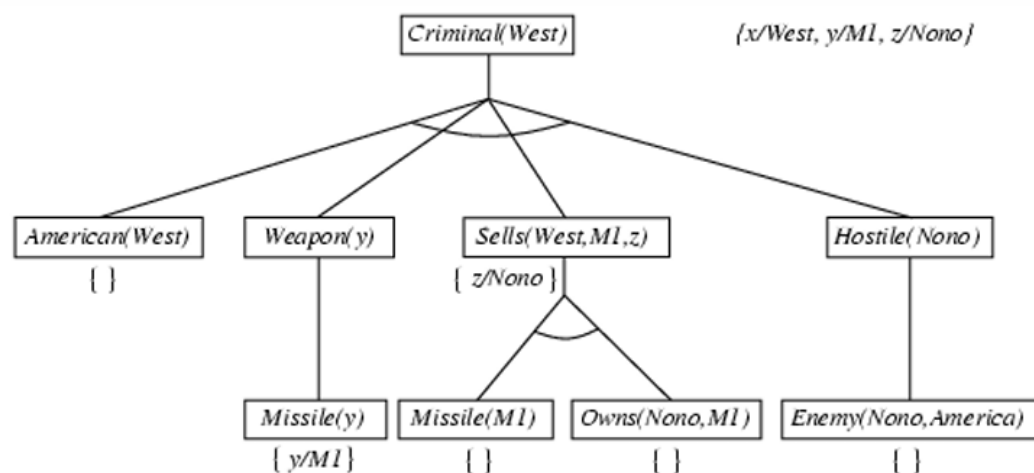
Scenario: Prove that *Colonel West is a criminal*.

1. **Goal:** Determine if Colonel West is a criminal.
2. **Facts:**
 - Colonel West is American.
 - Nono is a hostile nation.
 - Missiles are weapons.
3. **Rules:**
 - If an American sells weapons to a hostile nation, then they are a criminal.
 - If someone is American and has sold missiles to Nono, they have sold weapons.

Process:

- Start with the goal: "Colonel West is a criminal".
- Check the rule: "If an American sells weapons to a hostile nation, then they are a criminal".
- Sub-goal: Prove "Colonel West sold weapons to a hostile nation".
- Sub-goal: Prove "Colonel West is American" (already a fact).
- Sub-goal: Prove "Colonel West sold missiles to Nono" (derived from facts and rules).
- Work backward until all conditions are satisfied.

Backward chaining example



Comparison Between Forward and Backward Chaining

Feature	Forward Chaining	Backward Chaining
Type of Reasoning	Data-driven	Goal-driven
Starting Point	Known facts	Specific goal
Application	Generates all possible conclusions	Focuses only on proving the goal
Efficiency	May generate irrelevant conclusions	Focused and goal-specific
Example Use Cases	Real-time systems, expert systems	Diagnosis systems, problem-solving

STRIPS Representation for Block World Problem

Definition:

STRIPS (Stanford Research Institute Problem Solver) is a planning language used to represent states, actions, and goals in an environment. It is particularly useful for automated planning problems. In the context of the *block world problem*, STRIPS helps to formalize how blocks can be manipulated using operators like **pickup**, **putdown**, **stack**, and **unstack**.

Components of STRIPS Representation**1. States:**

A state is described using a set of predicates that define the configuration of the world. For example:

- $\text{On}(A, B) \rightarrow$ Block A is on block B.
- $\text{OnTable}(A) \rightarrow$ Block A is directly on the table.
- $\text{Clear}(A) \rightarrow$ Block A has no block on top of it.
- $\text{Holding}(A) \rightarrow$ The agent is holding block A.

2. Operators (Actions):

Each operator is defined using three components:

- **Preconditions:** The conditions that must hold true before the action can be executed.
- **Add List:** The predicates that become true after the action is executed.
- **Delete List:** The predicates that are no longer true after the action is executed.

3. Goals:

The desired final configuration of the blocks, expressed using predicates.

STRIPS Operators for Block World Problem**1. Pickup(X)**

- **Description:** Picks up block X from the table.
 - **Preconditions:**
 - $\text{OnTable}(X) \rightarrow$ Block X is on the table.
 - $\text{Clear}(X) \rightarrow$ Block X has no block on top.
 - $\text{HandEmpty} \rightarrow$ The agent's hand is free.
 - **Add List:**
 - $\text{Holding}(X) \rightarrow$ The agent is now holding block X.
 - **Delete List:**
 - $\text{OnTable}(X) \rightarrow$ Block X is no longer on the table.
 - $\text{HandEmpty} \rightarrow$ The agent's hand is no longer free.
-

2. Putdown(X)

- **Description:** Puts block X on the table.
 - **Preconditions:**
 - $\text{Holding}(X) \rightarrow$ The agent is holding block X.
 - **Add List:**
 - $\text{OnTable}(X) \rightarrow$ Block X is now on the table.
 - $\text{HandEmpty} \rightarrow$ The agent's hand is free.
 - **Delete List:**
 - $\text{Holding}(X) \rightarrow$ The agent is no longer holding block X.
-

3. Stack(X, Y)

- **Description:** Places block X on top of block Y.
 - **Preconditions:**
 - Holding(X) → The agent is holding block X.
 - Clear(Y) → Block Y has no block on top.
 - **Add List:**
 - On(X, Y) → Block X is now on block Y.
 - HandEmpty → The agent's hand is free.
 - **Delete List:**
 - Holding(X) → The agent is no longer holding block X.
 - Clear(Y) → Block Y is no longer clear.
-

4. Unstack(X, Y)

- **Description:** Removes block X from block Y.
 - **Preconditions:**
 - On(X, Y) → Block X is on block Y.
 - Clear(X) → Block X has no block on top.
 - HandEmpty → The agent's hand is free.
 - **Add List:**
 - Holding(X) → The agent is now holding block X.
 - Clear(Y) → Block Y is now clear.
 - **Delete List:**
 - On(X, Y) → Block X is no longer on block Y.
 - HandEmpty → The agent's hand is no longer free.
-

Example

Problem Description:

There are three blocks: A, B, and C. Initially:

- Block A is on Block B.
- Block B is on the table.
- Block C is on the table.
- The goal is to achieve the configuration where:
 - Block A is on Block C.
 - Block B remains on the table.

STRIPS Representation:

1. Initial State:

On(A,B),OnTable(B),OnTable(C),Clear(A),Clear(C),HandEmptyOn(A, B), OnTable(B), OnTable(C), Clear(A), Clear(C), HandEmptyOn(A,B),OnTable(B),OnTable(C),Clear(A),Clear(C),HandEmpty

2. Goal State:

On(A,C),OnTable(B)On(A, C), OnTable(B)On(A,C),OnTable(B)

3. Plan:

- **Unstack(A, B):** Preconditions: On(A, B), Clear(A), HandEmpty.
 - **Putdown(A):** Preconditions: Holding(A).
 - **Pickup(C):** Preconditions: OnTable(C), Clear(C), HandEmpty.
 - **Stack(A, C):** Preconditions: Holding(A), Clear(C).
-

Advantages of STRIPS Representation

- Provides a formal and systematic way to define planning problems.
- Efficient for solving problems in robotics and AI planning.
- Easily extendable to more complex domains.

Note: STRIPS is widely used in AI due to its simplicity and compatibility with modern planners.

MODULE 4 Graph Planning - Google Drive

Graph Planning in AI

Graph Planning is a technique that uses a planning graph to represent actions and states in solving a problem. The graph consists of alternating **state levels** and **action levels**, where each action is linked to its preconditions and effects. It helps in providing more accurate heuristics for planning compared to simpler methods.

Key Concepts:

1. **Planning Graph Structure:**
 - **State Level:** Contains literals (propositions) that can be true at a given time.
 - **Action Level:** Contains actions whose preconditions are satisfied by the literals in the previous state level.
2. **Persistence Actions:**
 - Represent inaction or the persistence of a literal when no action changes it (e.g., Have(Cake) remains true until an action like Eat(Cake) changes it).
3. **Mutual Exclusion (Mutex):**
 - Mutex links represent actions that cannot occur together due to conflicts:
 - **Inconsistent Effects:** Actions with conflicting outcomes (e.g., Eat(Cake) and Have(Cake)).
 - **Interference:** One action's effect negates another's precondition (e.g., Eat(Cake) negates Have(Cake)).
 - **Competing Needs:** Actions that compete for the same precondition (e.g., Bake(Cake) and Eat(Cake)).
4. **Leveled Off:**
 - The graph "levels off" when two consecutive levels are identical, meaning no new information can be added, and further expansion is unnecessary.

The GRAPHPLAN Algorithm:

1. **Expansion:**
 - The algorithm starts with the initial state and expands the graph by adding actions and state literals based on preconditions and effects.
2. **Goal Checking:**
 - It checks if all goal literals are present in the current state and ensures no mutex links between them. If true, it tries to extract a solution.
3. **Solution Extraction:**
 - If no solution is found, the graph is expanded further. The process repeats until a solution is found or deemed impossible.

Example: Spare Tire Problem

1. **Initialization:**
 - The graph starts with the initial state (e.g., At(Spare, Trunk)), and the goal (e.g., At(Spare, Axle)) is not present initially.
2. **Expansion:**
 - Actions whose preconditions are satisfied are added (e.g., Remove(Spare, Trunk)). Mutex relations (e.g., Remove(Spare, Trunk) and LeaveOvernight) are identified.
3. **Mutex Relations:**
 - Mutex relations include:
 - **Inconsistent Effects:** Conflicting effects of actions.
 - **Interference:** One action negates another's precondition.

- **Competing Needs:** Actions competing for the same precondition.

Advantages:

1. **Improved Heuristics:** More accurate than total-order or partial-order planning.
2. **Efficiency:** Avoids exploring infeasible solutions by using mutex constraints.
3. **Solution Extraction:** Provides a systematic way to extract solutions.

Disadvantages:

1. **Complexity:** Can become computationally expensive for large problems.
2. **Memory Usage:** Requires significant memory to maintain the planning graph and mutex relations.

Difference the Procedural and Declarative Knowledge

PROCEDURAL KNOWLEDGE	DECLARATIVE KNOWLEDGE
It is also known as Interpretive knowledge.	It is also known as Descriptive knowledge.
Procedural Knowledge means how a particular thing can be accomplished	While Declarative Knowledge means basic knowledge about something.
Procedural Knowledge is generally not used means it is not more popular.	Declarative Knowledge is more popular.
Procedural Knowledge can't be easily communicate.	Declarative Knowledge can be easily communicate
Procedural Knowledge is generally process oriented in nature	Declarative Knowledge is data oriented in nature.
In Procedural Knowledge debugging and validation is not easy.	In Declarative Knowledge debugging and validation is easy.

Inference in Predicate Logic

Inference in predicate logic refers to the process of deriving new information from a set of premises using logical rules. The following are common methods and rules used for inference:

1) Modes Pollen Method

The **Modes Pollen method** is a type of rule-based reasoning used to identify and eliminate possibilities in a logical system. It's more applicable in domain-specific contexts, like automated reasoning or expert systems. The method typically involves using heuristics to guide the search for valid conclusions based on existing knowledge.

2) Modes Tollen Method

The **Modes Tollen method** is another technique used for reasoning in predicate logic, often applied in the analysis of relationships or properties in systems. Like the Pollen method, it helps reduce the search space by evaluating possible predicates systematically and eliminates those that are inconsistent with the known premises.

3) AND Elimination

AND Elimination (also called **Conjunction Elimination**) is a rule that states if you have a conjunction $P \wedge Q$, you can infer both P and Q individually. This is a basic rule in predicate logic, and it is written as:

- From $P \wedge Q$, infer P .
- From $P \wedge Q$, infer Q .

4) Chain Rule

The **Chain Rule** (also known as **Hypothetical Syllogism**) is used to derive a conclusion based on two implications. If $P \Rightarrow Q$ and $Q \Rightarrow R$ are true, then $P \Rightarrow R$ can be inferred. It's widely used in predicate logic for connecting statements through intermediate propositions.

- From $P \Rightarrow Q$ and $Q \Rightarrow R$, infer $P \Rightarrow R$.

5) Substitution Method

The **Substitution Method** involves replacing variables in logical formulas with specific terms or predicates. This method is central to many logical systems, including unification in predicate logic. If you have a predicate $P(x)$ and know $x = a$, you can substitute a into $P(x)$ to get $P(a)$.

6) Universal Instantiation

Universal Instantiation is a rule that allows you to derive a specific case from a universally quantified statement. If $\forall x P(x)$ is true (meaning "P(x) is true for all x"), you can infer that $P(a)$ is true for any specific a .

- From $\forall x P(x)$, infer $P(a)$.

7) Existential Instantiation

Existential Instantiation is the reverse of universal instantiation. It allows you to derive a specific individual from an existentially quantified statement. If $\exists x P(x)$ is true (meaning "there exists an x such that P(x) is true"), you can instantiate $P(x)$ for a specific term a .

- From $\exists x P(x)$, infer $P(a)$ for some a .

8) Unification and Lifting

Unification is the process of finding a substitution that makes two logical expressions identical. For example, in predicate logic, unifying $P(x)$ with $P(a)$ involves substituting a for x . **Lifting** refers to generalizing or applying a specific substitution to higher-level formulas, extending the application of a substitution.

9) Inference Rules

Inference rules are the fundamental mechanisms in logic that govern the reasoning process. They allow the derivation of new statements (or conclusions) from premises. Some common inference rules in predicate logic include:

Inference Rules

1. Idempotent rule

$$P \wedge P \Rightarrow P$$

$$P \vee P \Rightarrow P$$

2. Commutative rule

$$P \wedge Q \Rightarrow Q \wedge P$$

$$P \vee Q \Rightarrow Q \vee P$$

3. Associative rule:

$$P \wedge (Q \wedge R) \Rightarrow (P \wedge Q) \wedge R$$

$$P \vee (Q \vee R) \Rightarrow (P \vee Q) \vee R$$

4. Distributive rule:

$$P \vee (Q \wedge R) \Rightarrow (P \vee Q) \wedge (P \vee R)$$

$$P \wedge (Q \vee R) \Rightarrow (P \wedge Q) \vee (P \wedge R)$$

5. De-Morgan's Rule:

$$\neg (P \vee Q) \Rightarrow \neg P \wedge \neg Q$$

$$\neg (P \wedge Q) \Rightarrow \neg P \vee \neg Q$$

6. Implication elimination:

$$P \rightarrow Q \Rightarrow \neg P \vee Q$$

Conclusion

Inference in predicate logic is an essential aspect of automated reasoning and logical deduction. By applying various methods like universal instantiation, substitution, and rules such as **AND Elimination** and **Chain Rule**, complex logical conclusions can be derived from simpler premises.